An LLM-driven multi-agent simulation framework for quantitative analysis of narrative voice shifts in literary texts

Lili Zhao[1] **and Hui Shi**[2]

[1]**Department of Foreign Languages, Bengbu University, Bengbu, Anhui, China, 233030**
[2]**Music College, Shanghai Normal University, Shanghai, China**

Corresponding author:
Lili Zhao[1]

Email address: jesszhaoll@163.com

ABSTRACT

Narratological accounts of first- and third-person narration rely mainly on interpretation and are difficult to test under controlled conditions. We present a large language model (LLM)-driven multi-agent simulation framework that makes narrative voice an experimental variable: a Director Agent extracts the characters of a literary excerpt and enforces the mode in which every turn is rendered, while Actor Agents with a sliding-window memory generate the dialogue. We evaluated the framework on six excerpts by six authors with three open-weight LLMs and five random seeds, using paired tests, a no-instruction control, a rule-based transposition baseline, ablations, two LLM judges and a small expert evaluation. The framework controlled narrative voice reliably (0.0% residual mode violations). In our simulations, third-person narration used longer and less frequent content words than first-person speech ($|d_z| = 0.90$) and was affectively more neutral ($|d_z| = 0.45$), effects that pronoun substitution does not explain, whereas lexicon-based subjectivity did not differ ($|d_z| = 0.03$). Character identity was only partially preserved across modes (speaker attribution 0.63 across vs. 0.84 within a mode; chance 0.38). A speak-then-narrate variant, in which the Director narrates the Actors' own speech in free indirect style, raised cross-mode attribution to 0.73 and removed most of the register shift, but the judges and the experts rated it as less coherent. These results support multi-agent LLM simulations as controllable instruments for computational narratology and show that single-run effects on one text do not necessarily generalize.

Project page: https://yanlibengbu.github.io/narrative-voice-sim/
Code and data: https://github.com/YanLiBengbu/narrative-voice-sim

INTRODUCTION

The analysis of narrative perspective has been a core theme in narratology since its formalization in structuralist theory. Shifts between first-person and third-person narration shape how readers perceive who speaks, whose perceptions are rendered, and how much a character's subjectivity is foregrounded (Onega and Landa, 1996; Bal, 2009). Classical narratologists emphasized close reading and interpretive frameworks (Baroni and Paschoud, 2021). Such approaches yield rich interpretations but are difficult to replicate and to test empirically. With the advent of computational methods, literary scholars have turned toward quantitative analysis, seeking systematic evidence for narratological phenomena (Gius and Vauth, 2022; Hatavara et al., 2024; Piper et al., 2021).

In digital literary studies, computational approaches have been applied to authorship attribution (He et al., 2024), stylistic analysis (Herrmann et al., 2021; Khafaga and Shaalan, 2022), the detection of narrative point of view and diegesis (Eisenberg and Finlayson, 2016), the quantification of free indirect discourse (Brooke et al., 2017), and quotation attribution in novels (Vishnubhotla et al., 2022). Most of these methods classify or describe existing texts; they leave little room for controlled experiments in which a single narratological variable is varied while everything else is held constant. Large language

models (LLMs) now make such experiments feasible, because they can extract character attributes from a passage and generate role-specific dialogue (Yu et al., 2025; Wang et al., 2024; Zhang et al., 2025a; Shao et al., 2023).

LLMs have also been used as interacting agents that take on distinct roles (Li et al., 2024; Sun et al., 2025; De Zarzà et al., 2023; Park et al., 2023), for example to model social dynamics, collaborative planning and conversational reasoning (Wang, 2025; Rahimi et al., 2025; Zhai et al., 2025). However, computational narratology has rarely combined multi-agent simulation with explicit control over narrative voice (Cai et al., 2024; Zhang et al., 2025b), so it remains open how the narrative form in which a scene is rendered changes the affective and stylistic profile of the resulting text.

In this work, we propose an LLM-driven multi-agent simulation framework for studying narrative voice shifts under controlled conditions. A Director Agent extracts the characters and dialogue of a literary excerpt and enforces the narrative mode of every turn: first-person direct speech, third-person narration by an external narrator, or a control condition without any perspective instruction. Actor Agents with a bounded dialogue memory generate the turns, and a seeded speaker-selection policy decides who speaks next. The system does not reproduce the source novel: it generates new dialogue conditioned on a short excerpt, and all results describe these generated dialogues. Our findings therefore characterize how LLM agents realize narrative-voice instructions; they can support, but not establish, narratological hypotheses about human-authored narrative.

We evaluate the framework on six excerpts by six authors with three open-weight LLMs and five random seeds. The main contributions are as follows:

- A multi-agent simulation framework in which a Director Agent controls narrative voice while Actor Agents with a sliding-window memory and an addressee-aware speaker-selection policy generate role-specific dialogue; the code, prompts, generated dialogues and analysis scripts are released.
- A controlled evaluation protocol: paired comparisons across texts, models and seeds, a no-instruction control and a rule-based transposition baseline that separate the direct effect of the mode instruction from changes that go beyond pronoun substitution, together with ablations of the main components.
- A test of character identity across narrative modes that calibrates a distributional consistency measure against a between-character baseline and complements it with speaker attribution, LLM judges and a small expert evaluation.

1 RELATED WORK

1.1 Narrative voice, perspective and focalization

Narratology distinguishes between two questions that the everyday term "point of view" tends to merge. *Voice* concerns who narrates: Genette describes the narrator's relation to the story as homodiegetic, when the narrator is a character in the story, or heterodiegetic, when the narrator is absent from it (Genette, 1980). *Focalization* concerns whose perception orients the narrative, independently of who narrates; it may be zero (unrestricted), internal (restricted to a character's perceptions) or external (Genette, 1980; Bal, 2009; Niederhoff, 2013). *Narrative perspective* is often used as an umbrella term for both dimensions (Niederhoff, 2013; Fludernik, 2009). Grammatical person is the most visible marker of voice, but it does not determine focalization: a third-person narrative can be internally focalized through a character, as in free indirect discourse (Brooke et al., 2017).

In this paper, the experimental variable is narrative voice in a narrow sense: the grammatical person and the narratorial mediation with which a character's turn is rendered. In the first-person mode, the turn is the character's own direct speech; in the third-person mode, a heterodiegetic narrator reports it. Focalization is not manipulated directly, although the narrator may mention the character's actions or inner states. "Perspective" below refers to this voice manipulation.

1.2 Empirical and computational studies of narrative perspective

Reader-response experiments show that grammatical person affects how narratives are processed. Pronouns modulate the perspective that readers adopt when they mentally simulate described events (Brunyé et al., 2009), and literary short stories presented with first-person pronouns produced higher immersion, mainly through transportation and mental imagery, than the same stories with third-person pronouns,

whereas appreciation was unaffected (Hartung et al., 2016). On the computational side, point of view and diegesis can be identified automatically with high accuracy, chiefly from pronoun features (Eisenberg and Finlayson, 2016), lexical-style models have been used to quantify free indirect discourse (Brooke et al., 2017), and annotated corpora support quotation and speaker attribution in novels (Vishnubhotla et al., 2022). Piper et al. (2021) argue that computational work on narrative should be connected more closely with narrative theory. These studies observe or classify existing texts; they do not vary narrative voice while holding the content of a scene constant.

1.3 Agent-Based Modeling in Narrative Contexts

The use of computational agents to simulate narrative dynamics has roots in interactive storytelling and computational narratology. Early systems such as Façade (Mateas and Stern, 2003) modeled characters as autonomous agents whose behavior responds to player interaction in real time. Cavazza et al. (2002) and Si et al. (2005) extended this paradigm by developing multi-agent systems where character goals dynamically shape story progression. Riedl and Young (2010) proposed narrative planning frameworks that balance plot coherence with character autonomy, while Swartjes and Theune (2006) introduced emergent narrative models where stories arise organically from agent interactions rather than pre-scripted sequences. Porteous and Cavazza (2009) further refined control mechanisms through planning trajectories, and Magerko et al. (2009) studied cognition in theatrical improvisation to inform agent-based models of improvised performance. However, these traditional systems relied on hand-crafted rules and symbolic planning, limiting linguistic flexibility and scalability. Our work keeps the principle of simulating narrative through autonomous agents but uses LLMs to generate the dialogue and to control its narrative voice.

1.4 LLM-based Role Play and Multi-Agent Systems

Recent advances in large language models (LLMs) have enabled the development of multi-agent frameworks, where autonomous agents interact to simulate complex social and linguistic dynamics. Generative agents with memory and reflection can simulate believable everyday behavior (Park et al., 2023), and role play has been proposed as a framework for understanding what dialogue agents do (Shanahan et al., 2023). Dedicated methods train or evaluate LLMs for character role play and plot progression (Shao et al., 2023; Yu et al., 2025; Wang et al., 2024; Zhang et al., 2025a). Guo et al. (2024) provide a comprehensive survey of LLM-based multi-agent systems that identifies open issues such as benchmarking. LLMs are also increasingly used as evaluators of generated text; their ratings correlate with expert judgements in several tasks but are subject to biases such as position and self-preference effects (Zheng et al., 2023; Chiang and Lee, 2023), so we use LLM judges only as a complement to automatic measures and human ratings. Our framework positions multi-agent LLM systems not as content generators but as experimental instruments for the controlled analysis of narrative voice.

2 METHODOLOGY

2.1 System Framework

The proposed system adopts an LLM-driven multi-agent framework to simulate and analyze narrative voice shifts in literary texts. As illustrated in Fig. 1, the framework is organized into four main modules: role extraction, dialogue simulation, dialogue analysis, and comparative reporting. The system starts by taking a literary excerpt as input and initializes the environment through a *Director Agent*, which is responsible for character assignment, context extraction, and narrative mode configuration. Based on the extracted speaker–listener pairs and character traits, a set of *Actor Agents* is instantiated, each bound to a specific literary character. The excerpt is used solely to initialize agents and context; all subsequent analysis is conducted on the newly generated dialogues, not on the source text. The generated dialogues are then processed by the *Dialogue Analyzers*, which compute utterance- and session-level measures, and by a *Vocabulary Consistency Analyzer*, which compares character-specific lexical distributions across modes. Finally, results are aggregated into mode-wise summaries, statistical tests and reports.

2.2 Dialogue Simulation and Narrative Modes

Scene initialization. The Director Agent first extracts the direct-speech turns of the excerpt as $\{$speaker, addressees, text$\}$ records and then writes a role card for every character (name, gender, personality, emotional baseline in $[-1, 1]$, speech style, background and relationships) and a short scene context. Because LLM speaker labels are often inconsistent (e.g., “she”, “the lady” and “Miss Havisham” for the

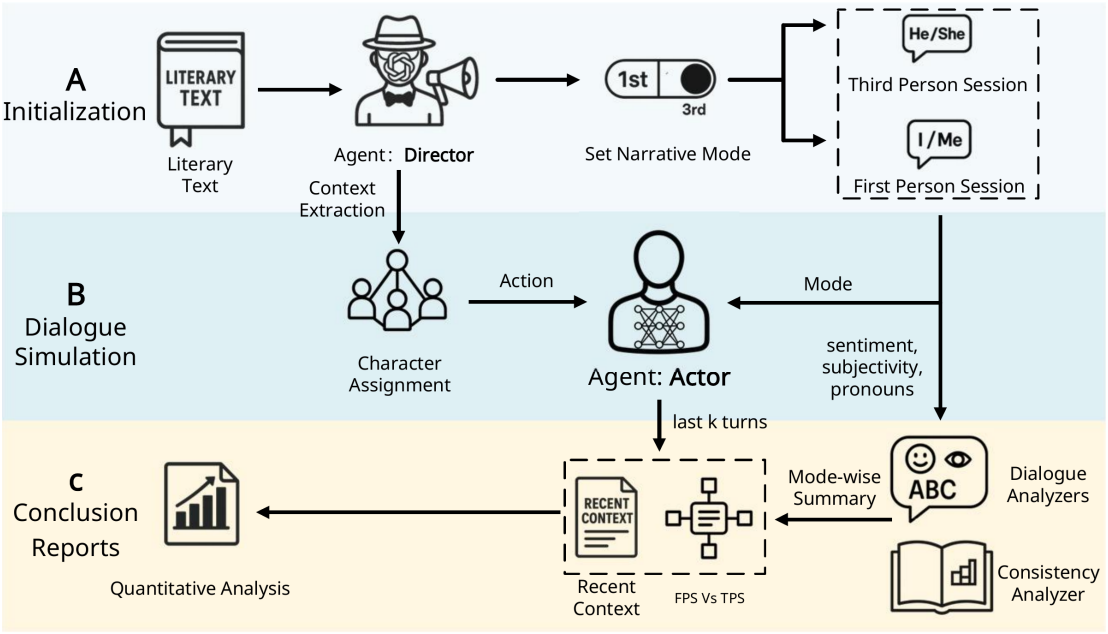


Figure 1. Overall architecture of the proposed LLM-driven multi-agent simulation framework. The system consists of role extraction, dialogue simulation under different narrative modes, multi-dimensional dialogue analysis, and comparative reporting.

same person), the labels are mapped to a manually verified cast list with aliases; turns whose speaker cannot be mapped are discarded. Scene initialization is performed once per excerpt and model with temperature 0, and all modes and seeds share the same role cards and opening turns, so that the narrative mode is the only factor that differs between paired sessions.

Narrative modes. In the *first-person* mode, each turn is rendered as the character's own direct speech, without narratorial framing. In the *third-person* mode, an external (heterodiegetic) narrator reports what the character says, may mention actions or inner states, refers to the character by name or by a third-person pronoun, and uses no quotation marks; the opening turns of the excerpt are rewritten into reported narration by the Director. The *neutral* control contains no perspective instruction. The instructions are given verbatim in the Appendix.

Mode enforcement. After each turn, the Director applies a rule-based compliance check. A third-person turn violates the mode if it contains quotation marks or first- or second-person pronouns, or if it does not refer to the speaker by name or a third-person pronoun; a first-person turn violates the mode if it describes the speaker in the third person with a reporting verb. A non-compliant turn is regenerated by the Actor with an explicit hint and, if it still fails, rewritten by the Director. For every turn we record whether it complied at the first attempt, was regenerated, was rewritten, or remained non-compliant.

Actor Agents and memory. Each Actor Agent receives its role card, the scene context, the names of the other characters, the last k turns of the dialogue (sliding-window memory, $k = 5$ by default) and the mode instruction, and returns a JSON object containing the intention (*speak* or *silent*), the content of the turn and its addressees.

Speaker selection. The next speaker is sampled from a softmax over addressee-aware scores computed from the truncated history H'. For character c_i with g_i turns since its last turn,

$$s_i = \alpha \mathbb{1}[c_i \in A_{t-1}] - \beta \mathbb{1}[c_i = c_{t-1}] + \gamma \frac{\min(g_i, |\mathscr{C}|)}{|\mathscr{C}|}, \qquad p_i = \frac{\exp(s_i/\tau)}{\sum_j \exp(s_j/\tau)}, \tag{1}$$

where A_{t-1} and c_{t-1} are the addressees and the speaker of the previous turn, $\alpha = 2.0$, $\beta = 1.5$, $\gamma = 1.0$ and $\tau = 0.7$. The scores favor characters who were just addressed, penalize immediate self-succession, and give characters who have not spoken recently a chance to take the floor; sampling is seeded with the session seed.

Algorithm 1 Multi-agent dialogue generation with narrative-mode enforcement and reproducibility control

Input: scene (role cards, context, opening turns H_0), mode m, seed s, memory k, maximum turns $T_{\max}$
Output: dialogue $D = \{d_1, \ldots, d_T\}$
Initialize RNG(s), $H \leftarrow H_0$, $D \leftarrow \emptyset$
while $|D| < T_{\max}$ **and** fewer than 4 consecutive silent turns **do**
  $H' \leftarrow$ last k turns of H {sliding-window memory}
  compute s_i for all $c_i \in \mathscr{C}$ (Eq. 1); $c^* \sim \text{softmax}(s/\tau)$
  $r \leftarrow \text{Actor}_{c^*}(\text{card}_{c^*}, \text{context}, H', m;\ \text{seed } s_t)$ {JSON output}
  if r is a spoken turn **and** violates mode m **then**
    regenerate with a hint; if still violating, Director rewrites r into mode m
  end if
  $D \leftarrow D \cup \{r\}$; $H \leftarrow H \cup \{r\}$
end while
return D

Algorithm 1 summarizes the generation procedure. Every LLM call uses a seed derived from the session seed and the turn index. JSON outputs are requested with constrained decoding and parsed with a deterministic repair routine under a budget of three attempts. All raw LLM responses are cached, so that every reported number can be regenerated from the released outputs.

2.3 Quantitative Analysis Methods

Figure 2 shows the analysis pipeline.

Manipulation checks. The share of first-person pronouns among first- and third-person pronouns (*first-person ratio*) and the number of self-references by name are direct consequences of the mode instruction; they verify the manipulation and are not treated as findings.

Outcome measures. Sentiment is measured at the utterance level with a DistilBERT classifier fine-tuned on SST-2 (Sanh et al., 2019) (signed score $p_{\text{pos}} - p_{\text{neg}}$), with TextBlob polarity (Loria, 2020) and with VADER (Hutto and Gilbert, 2014); *emotional intensity* is the absolute TextBlob polarity. Subjectivity is the TextBlob subjectivity score. Utterance length is the number of word tokens. To prevent the trivial substitution of "I" by a character name from driving the lexical measures, word length and word frequency are computed on *content words* only, i.e., alphabetic tokens that are not stop words, pronouns, proper nouns or parts of character names (spaCy tokenization and tagging; Honnibal et al., 2020). Word frequency is the mean Zipf frequency of the content words according to the `wordfreq` lexicon (Speer, 2022), and the *rare-word rate* is the share of content words with Zipf frequency below 3. Lexical diversity is measured with the length-robust moving-average type–token ratio (MATTR, window 50; Covington and McFall, 2010) and the measure of textual lexical diversity (MTLD; McCarthy and Jarvis, 2010).

Interaction measures. Turn-taking balance is the normalized entropy of the distribution of turns over characters (1 = perfectly balanced); addressee responsiveness is the share of turns whose speaker was addressed in the previous turn; self-succession is the share of consecutive turns by the same speaker; and repetition is one minus the proportion of distinct word trigrams in a session.

Character consistency. For a character c and mode m, we pool the content words and utterance-level sentiment classes over all seeds and compute the distributions of word length P_{len}, Zipf-frequency band P_{freq} and sentiment class P_{sent} (positive, neutral, negative). The Jensen–Shannon divergence (Lin, 1991)

$$\text{JS}(P \,\|\, Q) = \tfrac{1}{2}\text{KL}(P \,\|\, M) + \tfrac{1}{2}\text{KL}(Q \,\|\, M), \qquad M = \tfrac{1}{2}(P+Q), \tag{2}$$

with base-2 logarithms lies in $[0,1]$, and the composite consistency score is

$$S(P,Q) = 1 - \left(0.3\,\text{JS}_{\text{len}} + 0.4\,\text{JS}_{\text{freq}} + 0.3\,\text{JS}_{\text{sent}}\right). \tag{3}$$

Both distributions are subsampled to the same size (120 content words and 8 utterances, 20 draws). A character's *within-character* consistency across modes, $S(P_{c,\text{first}}, P_{c,\text{third}})$, is compared with its *between-character* consistency with the other characters of the scene within a mode; traits count as preserved only if the former exceeds the latter.

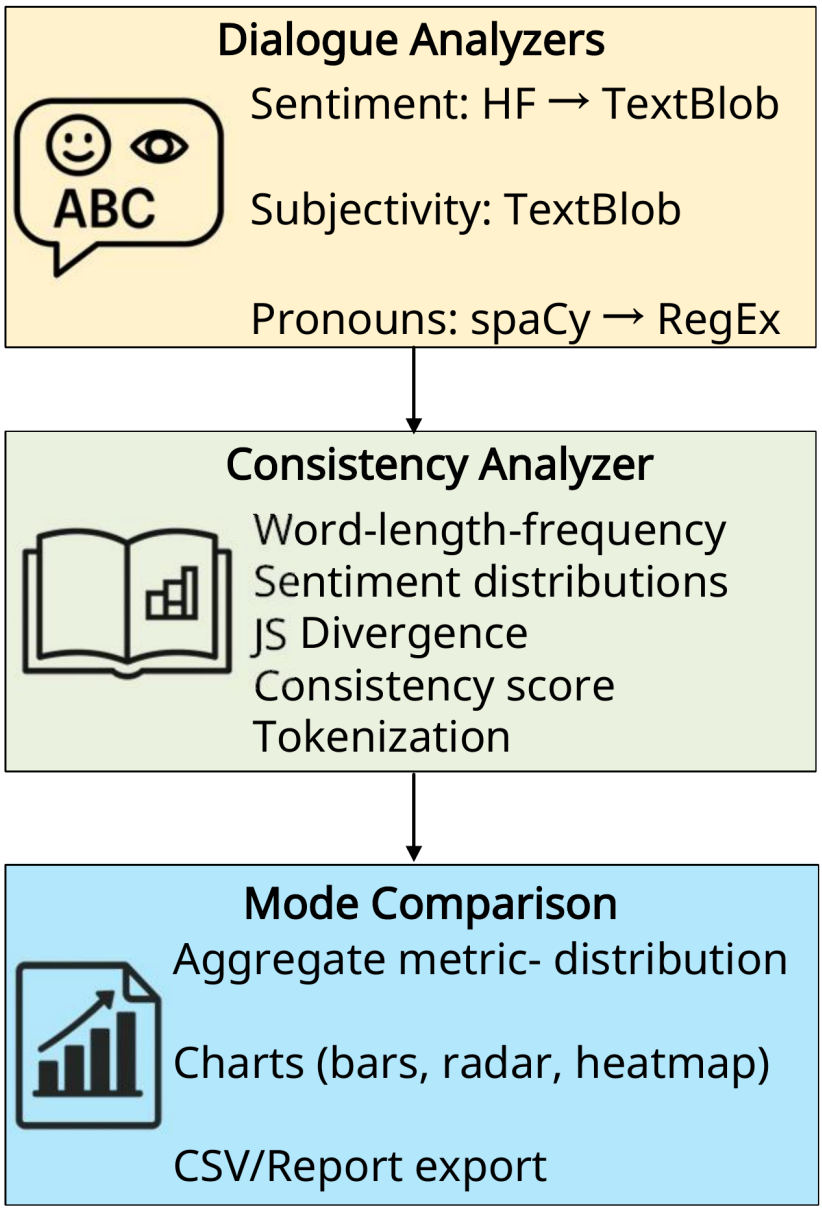


Figure 2. Software-level analysis pipeline. Dialogue analyzers capture sentiment, subjectivity, and pronoun statistics, while the consistency analyzer evaluates distributional features and outputs comparative metrics.

Speaker attribution. As a more sensitive test of character identity, we train, separately for every excerpt–model cell, a logistic-regression classifier that predicts the speaker of a turn from TF-IDF weighted content-word unigrams and bigrams (pronouns and names removed) or from the frequencies of 96 non-pronoun function words; characters with fewer than ten turns are omitted. Within-mode accuracy is estimated with leave-one-seed-out cross-validation, and cross-mode accuracy by training on one mode and testing on the other. We compare balanced accuracy with chance ($1/K$ for K characters) and with the within-mode accuracy across the 18 cells (Wilcoxon signed-rank tests).

2.4 Statistical procedures

The unit of analysis is the session: utterance-level values are first averaged within each session to avoid pseudo-replication. Sessions that share the excerpt, the model and the seed form a block, and modes are compared within blocks. For each measure and pair of modes we report the mean paired difference with a 95% bootstrap confidence interval (10,000 resamples of blocks; Efron and Tibshirani, 1993), the two-sided Wilcoxon signed-rank test (Wilcoxon, 1945), Cohen's d_z and the matched-pairs rank-biserial correlation r_{rb}. P-values are adjusted with the Holm procedure within each family of comparisons (Holm, 1979). As a robustness check, we fit linear mixed models with mode and generator model as fixed effects and a random intercept for the excerpt. To assess generalization, we also report in how many of the 18 excerpt–model cells the mean difference has the same sign. All tests are two-sided with $\alpha = 0.05$.

2.5 LLM-judge evaluation

As a scalable complement to the automatic measures, two judge models evaluated the generated dialogues at temperature 0: `phi4:14b` (Abdin et al., 2024), from a developer different from those of the three generators, and `qwen3.5:9b` (Qwen Team, 2026), from the same model family as one generator. Each judge rated every session on 1–5 scales for coherence, character consistency, scene fidelity, subjectivity and emotional intensity and classified the realized narrative form; it saw the excerpt metadata, the role cards and the dialogue, but not the generating model or the condition. In a speaker-identification task, four sampled turns per session (seeds 0–2) were shown with all names masked and gendered pronouns neutralized, and the judge identified the speaker among the scene's characters and rated the turn's subjectivity. Agreement between judges is quantified with Krippendorff's α (Hayes and Krippendorff, 2007). A blinded expert evaluation of a subset is described in subsection 3.10.

2.6 Computing infrastructure

All experiments were run on a workstation with Ubuntu 24.04.5 LTS, an NVIDIA GeForce RTX 4070 Ti SUPER GPU (16 GB), 24 CPU threads and 30 GB of RAM. LLMs were served locally with Ollama 0.40.3 (Ollama, 2026) using 4-bit quantized weights. The analysis used Python 3.12 with spaCy 3.8 (`en_core_web_sm`), TextBlob 0.20, NLTK 3.10 (VADER), Transformers 5.17, wordfreq 3, SciPy 1.17, statsmodels 0.15, scikit-learn 1.9 and krippendorff 0.9.

3 EXPERIMENTS

3.1 Experimental Setup

Texts. We use six dialogue-rich excerpts by six authors that differ in genre, period and narrative structure (Table 1): a scene from *The Good Earth* (Buck, 1931), Darcy's first proposal from *Pride and Prejudice* (Austen, 1813), Pip's first meeting with Miss Havisham from *Great Expectations* (Dickens, 1861), the visit of the masked King in "A Scandal in Bohemia" (Doyle, 1892), Lady Bracknell's interview of Jack in *The Importance of Being Earnest* (Wilde, 1899), and the confrontation between Nikhil and Sandip in *The Home and the World* (Tagore, 1919). They cover heterodiegetic narration, free indirect discourse, autodiegetic and homodiegetic first-person narrators, the dramatic mode and alternating narrators, and two of them have non-Western settings. The five additional texts are in the public domain (Project Gutenberg). For every excerpt we compiled the cast with name aliases and attributed every quotation to its speaker manually, to verify the cast and to evaluate the Director's extraction.

Models and parameters. Three open-weight instruction-tuned LLMs from different developers serve as generators: Llama 3.1 8B (Llama Team, AI @ Meta, 2024), Qwen2.5 14B (Qwen Team, 2024) and Gemma 3 12B (Gemma Team, 2025). All three are served locally by Ollama with 4-bit quantization. The same model plays the Director and all Actors within a session. Table 2 lists the configuration.

Table 1. Literary excerpts used in the experiments. Cast: characters represented by Actor Agents; Quotations: manually attributed direct-speech segments.

Excerpt	Source	Genre	Narrative structure	Words	Cast	Quotations
The Good Earth	Buck (1931), ch. 26	Realist family novel	Heterodiegetic, third person	901	4	11
Pride and Prejudice	Austen (1813), ch. 34	Novel of manners	Heterodiegetic, free indirect discourse	1063	2	14
Great Expectations	Dickens (1861), ch. 8	Bildungsroman	Autodiegetic, retrospective first person	531	2	24
A Scandal in Bohemia	Doyle (1891), part I	Detective story	Homodiegetic witness narrator	657	3	20
The Importance of Being Earnest	Wilde (1895), act I	Comic drama	No narrator (dramatic mode)	671	3	25
The Home and the World	Tagore (1916; tr. 1919), ch. 10	Political novel	Alternating first-person narrators	441	3	12

Table 2. Experimental parameters and configurations.

Parameter	**Value**
Generator LLMs	`llama3.1:8b`, `qwen2.5:14b`, `gemma3:12b` (Ollama 0.40.3, Q4_K_M)
Judge LLMs	`phi4:14b`, `qwen3.5:9b-q4_K_M` (temperature 0)
Narrative modes	first person, third person, neutral (no instruction); improvement experiments: voice-preserving third person, speak-then-narrate
Maximum dialogue turns $T_{\max}$	20 generated turns per session
Generation temperature	0.7 (Director extraction and rewriting: 0)
Maximum tokens per response	1000 (extraction and single-prompt baseline: 3000–4000)
Context window	8192 tokens
Memory window k	5 turns (ablation: 1)
Speaker selection	addressee-aware softmax, $\alpha = 2.0$, $\beta = 1.5$, $\gamma = 1.0$, $\tau = 0.7$ (ablations: uniform, round robin)
Random seeds	0–4 (main experiment, improvement 1); 0–2 (ablations, extension, improvement 2)
JSON retry budget	3 attempts per call; mode enforcement: 1 regeneration + Director rewrite
Sentiment analysis	DistilBERT (SST-2), TextBlob, VADER
Subjectivity analysis	TextBlob
Pronoun and lexical analysis	spaCy `en_core_web_sm`, `wordfreq` Zipf frequencies

Design and number of runs. The main experiment simulated every combination of excerpt (6), generator (3), narrative mode (3) and seed (5) once, giving 270 sessions with at most 20 generated turns each. Sessions with the same excerpt, model and seed form a block that shares the scene initialization, so that modes are compared in a paired design. The rule-based transposition baseline converts every first-person session into third-person reported speech (pronoun substitution plus a reporting frame such as "O-lan told Wang Lung that ...") without changing the content. The ablations used Llama 3.1 8B on all six excerpts, both modes and seeds 0–2 (36 sessions per condition). The extension and improvement experiments used all three generators (voice exemplars: both modes, seeds 0–2, 108 sessions; voice-preserving instruction: seeds 0–4, 90 sessions; speak-then-narrate: seeds 0–2, 54 sessions). In total, 702 sessions with 13704 generated turns were analyzed; 24 sessions ended early because the selected characters remained silent four times in a row.

3.2 Scene initialization and pipeline reliability

Table 3 summarizes the Director's extraction and the mode enforcement. Averaged over excerpts and models, the Director recovered 96.7% of the manually attributed quotations, attributed 93.2% of them to the correct speaker after alias mapping, and produced 2.5% unmatched utterances; typical errors were merged or split quotations and misattributed short replies in multi-party scenes. Because the cast was verified and the opening turns are shared by all modes within a block, such errors cannot create mode differences. With JSON-constrained decoding, none of the 15772 LLM calls produced unparseable output. In the third-person mode, 95.8% of the turns complied at the first attempt, 3.9% were corrected by regeneration and 0.3% were rewritten by the Director, leaving 0.0% non-compliant turns (first-person mode: 99.8% compliant at the first attempt).

Table 3. Director extraction quality and narrative-mode enforcement. Extraction is evaluated against manually attributed quotations (mean over six excerpts); enforcement figures are shares of generated turns in the main experiment. FP: first person; TP: third person.

Model	Quote recall	Attribution acc.	Fabricated	TP compliant (1st)	TP regenerated	TP rewritten	Residual (TP)	Residual (FP)
Llama 3.1 8B	0.96	0.95	0.02	0.932	0.060	0.008	0.000	0.000
Qwen2.5 14B	0.94	0.93	0.00	0.998	0.002	0.000	0.000	0.000
Gemma 3 12B	1.00	0.92	0.05	0.945	0.055	0.000	0.000	0.002

3.3 An example of the generated dialogues

Table 4 shows the first eight turns of two sessions generated by Qwen2.5 14B for the scene from *The Good Earth* with the same role cards, opening and seed. The first-person session consists of direct speech ("My wife, do not fret"); in the third-person session, a narrator reports the same kinds of speech acts ("Wang Lung reassured O-lan that ...") and occasionally adds gestures or inner states. Both sessions develop the situation of the source scene, O-lan's wish to see her eldest son married before she dies, but neither reproduces the novel. Transcripts for all excerpts are given in the Supplemental Information (S10).

3.4 Manipulation check and effects of narrative mode

Table 5 reports the session-level means in the three instructed modes and in the rule-based transposition, with the paired first- versus third-person comparisons; Fig. 3 shows the standardized differences for all pairs of conditions.

Manipulation check. The instructions were realized as intended: first-person pronouns accounted for 0.846 of the first- and third-person pronouns in first-person sessions and for 0.000 in third-person sessions. Without any instruction, the models also produced mainly first-person direct speech (first-person ratio 0.800), so first-person speech is the default voice of role-playing LLMs, and the third-person instruction is the manipulation that departs from it.

Subjectivity and affect. TextBlob subjectivity did not differ between the modes (difference 0.003, 95% CI $[-0.019, 0.024]$, $p_{\text{Holm}} = 1.000$; positive in 9/18 excerpt–model cells). Emotional intensity was higher in first-person sessions (difference 0.034, CI [0.018, 0.050], $d_z = 0.45$, $p_{\text{Holm}} < 0.001$; positive in 15/18 cells), and so was TextBlob polarity ($p_{\text{Holm}} = 0.025$), whereas VADER showed no reliable difference ($p_{\text{Holm}} = 1.000$). The rule-based transposition left intensity and polarity unchanged, so these differences are not produced by the reporting form. The DistilBERT difference (0.161, $p_{\text{Holm}} < 0.001$) is

Table 4. Example of generated dialogue in the two narrative modes. First eight turns of the *The Good Earth* scene generated by Qwen2.5 14B (seed 0) with identical role cards and opening.

Mode	Turn	Speaker	Generated turn
First person	1	Wang Lung	My wife, do not fret. Our son will marry well and care for the fool as you have. The gods will see to it that all is well.
	2	O-lan	I know it is hard to let go, but you must look after them as I cannot. Ensure they are cared for and that the fool is well taken care of.
	3	Wang Lung	My dear, you must rest now. The eldest son will marry soon, and I will see to it that all your wishes are fulfilled. Do not worry, for I am here and will take care of everything.
	4	O-lan	I trust you, my husband, and I know you will do right by them. But promise me you will see to the fool's needs as I have.
	5	the eldest son	Mother, please rest assured. I will take care of everything as Father says. I will marry soon and look after the fool as you have.
	6	Wang Lung	My son, your words are a comfort. I will see to it that the wedding preparations are made and that your mother's last wishes are honored.
	7	the eldest son	Mother, please try to rest. Father and I will make sure everything is taken care of.
	8	O-lan	Thank you, my son. And Father, I know you will see to it. I only ask that you do not forget the fool's happiness above all.
Third person	1	Wang Lung	Wang Lung reassured O-lan that everything would be taken care of, that their son would be guided well and that she had no need to worry about him or the family. He insisted that he would see her wishes fulfilled and that her eldest son would have a good life.
	2	O-lan	O-lan asked Wang Lung if he had ensured that their son would be taught the ways of respect and duty towards his wife and her family, and if he had reminded him of the importance of providing for his wife and children.
	3	Wang Lung	Wang Lung assured O-lan that he would see to it that their son would be taught the ways of respect and duty, and that he would remind him of the importance of providing for his wife and her family, as she wished.
	4	O-lan	O-lan asked Wang Lung if he had spoken to the maid about the importance of her role in the family, and if he had reminded her of the duties she owed to her husband and his ancestors.
	5	the eldest son	The eldest son, feeling a mix of nervousness and devotion, told his mother that he understood the importance of respecting his wife and her family and would do his best to provide for them.
	6	the maid	The maid told O-lan that she had understood her duties and would serve the family with all her heart, ensuring that she respected her husband and his ancestors.
	7	O-lan	O-lan told Wang Lung and the others that she felt reassured by their words, but she asked Wang Lung to ensure that the arrangements for her son's marriage were completed swiftly, for her peace of mind.
	8	Wang Lung	Wang Lung told O-lan that he would personally oversee the arrangements for the eldest son's marriage, ensuring that everything was done swiftly and in accordance with tradition to bring her peace of mind.

partly caused by the reporting frame, which by itself lowered the score by 0.061. The affective effects were strongest for Gemma 3 12B (Supplemental Information, S3).

Lexical register. The largest and most consistent differences concerned the vocabulary. Content words were longer in third-person narration (7.05 vs. 6.67 characters; $d_z = -0.90$, $p_{\text{Holm}} < 0.001$), less frequent ($d_z = 0.78$, $p_{\text{Holm}} < 0.001$) and more often rare ($p_{\text{Holm}} = 0.002$). The transposition changed both measures in the opposite direction, so the effect is not an artefact of substituting names for pronouns; it reflects the reporting verbs and the more formal diction of narration (e.g., *reassured*, *insisted*, *expressed*). Third-person turns were also longer (41.2 vs. 34.2 words), mostly because of the reporting frame, which by itself added 5.8 words. Lexical diversity was higher in third-person sessions by MATTR ($p_{\text{Holm}} = 0.009$) but not by MTLD ($p_{\text{Holm}} = 1.000$), so this difference is only partly supported.

Neutral control. The neutral sessions did not differ from the first-person sessions in emotional intensity ($p_{\text{Holm}} = 0.738$) or content-word length ($p_{\text{Holm}} = 0.738$), whereas the third-person sessions differed from both. Narration by an external narrator therefore dampens affect and raises the lexical register, rather than first-person speech amplifying affect.

Table 5. Session-level means by narrative mode and paired first- vs. third-person comparisons (main experiment: 6 excerpts × 3 models × 5 seeds; $n = 90$ blocks). †Manipulation check. Rule-based: first-person sessions transposed to third person by pronoun substitution and a reporting frame. Δ: mean paired difference (first − third) with 95% bootstrap CI; d_z: Cohen's d_z; p_{Holm}: Wilcoxon signed-rank test, Holm-adjusted across the measures of the table; Cells > 0: excerpt–model cells (of 18) with a positive first − third difference.

Measure	First	Third	Neutral	Rule-based	Δ First−Third [95% CI]	d_z	p_{Holm}	Cells > 0
First-person ratio†	0.846	0.000	0.800	0.000	0.845 [0.820, 0.869]	7.17	<0.001	18/18
Self-reference by name†	0.3	2.0	0.3	1.9	−1.7 [−1.8, −1.5]	−2.23	<0.001	0/18
Subjectivity	0.486	0.483	0.499	0.486	0.003 [−0.019, 0.024]	0.03	1.000	9/18
Emotional intensity	0.205	0.171	0.218	0.205	0.034 [0.018, 0.050]	0.45	<0.001	15/18
Polarity (TextBlob)	0.121	0.092	0.124	0.121	0.028 [0.010, 0.046]	0.33	0.025	16/18
Polarity (VADER)	0.208	0.189	0.254	0.199	0.019 [−0.032, 0.069]	0.08	1.000	10/18
Polarity (DistilBERT)	0.223	0.062	0.284	0.162	0.161 [0.088, 0.235]	0.45	<0.001	12/18
Utterance length (words)	34.2	41.2	32.5	40.1	−7.0 [−9.1, −5.0]	−0.71	<0.001	5/18
Content-word length	6.669	7.054	6.719	6.450	−0.386 [−0.474, −0.298]	−0.90	<0.001	2/18
Content-word Zipf frequency	4.616	4.501	4.602	4.690	0.114 [0.084, 0.145]	0.78	<0.001	15/18
Rare-word rate	0.039	0.048	0.041	0.037	−0.009 [−0.014, −0.004]	−0.38	0.002	4/18
MATTR	0.843	0.870	0.860	0.805	−0.027 [−0.043, −0.012]	−0.36	0.009	7/18
MTLD	70.9	72.5	76.5	62.4	−1.6 [−5.7, 2.3]	−0.08	1.000	9/18

3.5 Generalization across texts and models

Figure 4 breaks the first- versus third-person differences down by excerpt and model. Content-word length was higher in third-person narration in 16/18 cells, Zipf frequency was higher in first-person turns in 15/18 cells, and emotional intensity and polarity were higher in first-person turns in 15/18 and 16/18 cells; the DistilBERT difference (12/18), MATTR (11/18) and subjectivity (9/18) were less consistent. Linear mixed models with a random intercept for the excerpt led to the same conclusions as the paired tests (Supplemental Information, S2). Several effects depended on the generator: Gemma 3 12B produced much longer third-person than first-person turns (difference −17.9 words), while the other two models differed by fewer than five words. Effects found for one model on one text should therefore not be generalized without replication.

3.6 Character-level Analysis

Figure 5 and Table 6 show the character-level distributions for *The Good Earth*. In both modes, the four characters used similar content-word lengths, and third-person narration shifted all of them towards longer and less frequent words; the differences between characters were small compared with the differences between modes.

Figure 6 shows session-level subjectivity by mode and a Bland–Altman analysis (Bland and Altman, 1986) of character-level subjectivity. The mean difference was close to zero (0.003) with wide limits of agreement: switching the mode changed a character's subjectivity in both directions, depending on the character, the text and the model, so single runs can suggest a difference in either direction.

Are character-specific traits preserved across modes? Small Jensen–Shannon divergences between a character's first- and third-person turns indicate preserved traits only if the characters differ from one

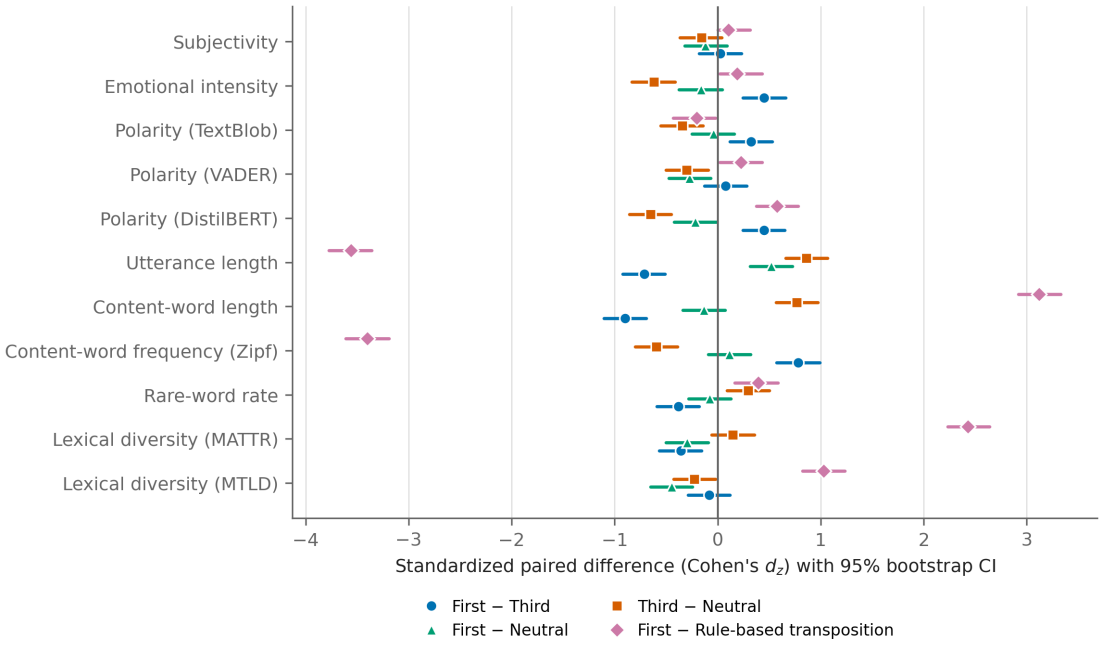


Figure 3. Standardized paired differences between narrative modes. Points show Cohen's d_z of the session-level differences for each measure and pair of conditions; bars show 95% bootstrap confidence intervals ($n = 90$ excerpt–model–seed blocks). Rule-based transposition: first-person sessions converted to third person by pronoun substitution and a reporting frame, without any change of content.

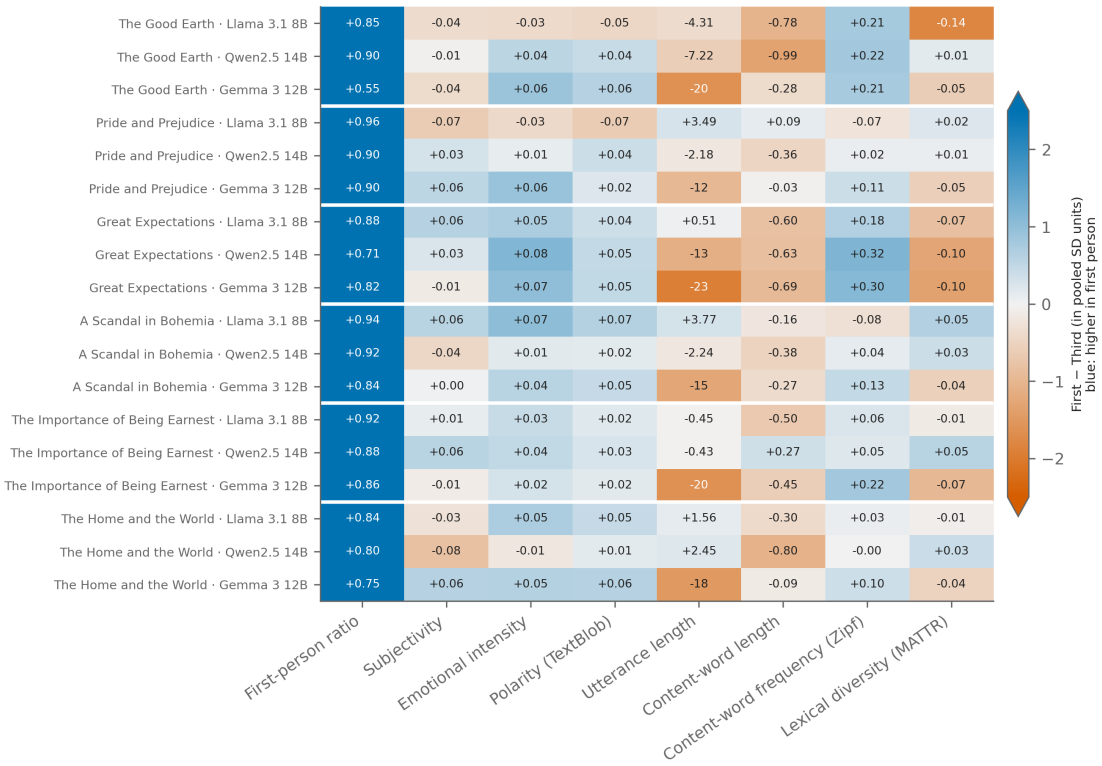


Figure 4. First- minus third-person differences for every excerpt and generator. Cells show the mean difference of the session-level measure over five seeds; colors show the difference in units of the pooled standard deviation (blue: higher in first person; orange: higher in third person; saturated at ± 2.5).

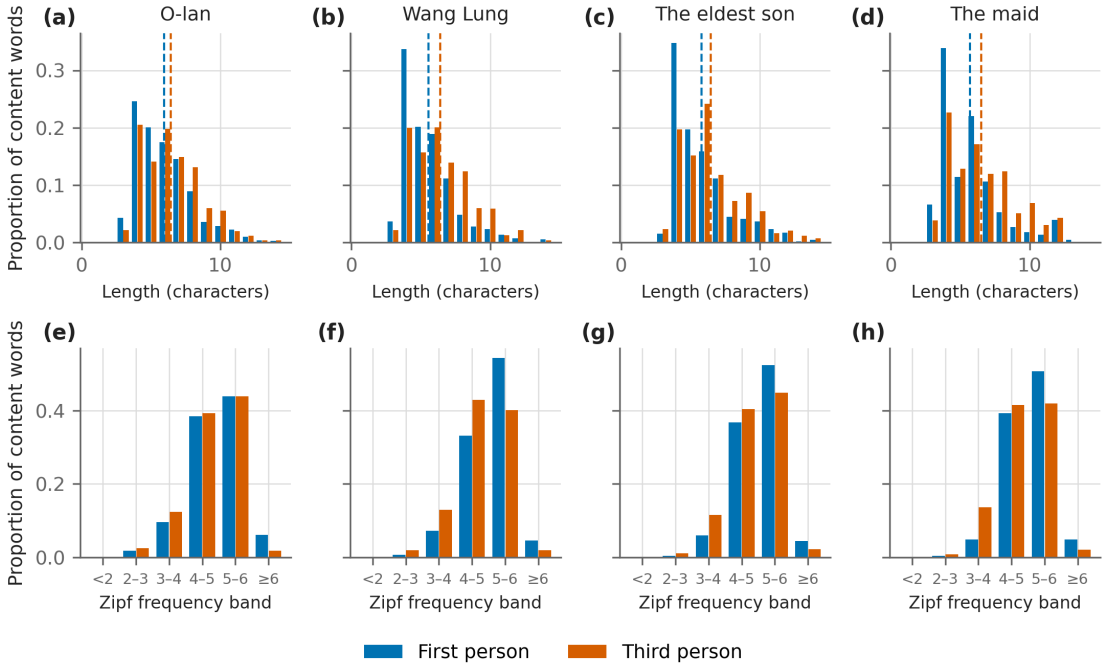


Figure 5. Content-word length (a–d) and word-frequency (e–h) distributions of the four characters of *The Good Earth* in first-person (blue) and third-person (orange) sessions (all models and seeds pooled). Dashed lines: mean word length. Frequency bands are Zipf frequencies from `wordfreq`.

Table 6. Character-level lexical statistics for *The Good Earth* (all models and seeds pooled). *S* within: composite consistency of the same character between its first- and third-person turns; *S* between: mean consistency with the other characters of the scene in the same mode (averaged over models; characters with fewer than 120 content words in a mode are omitted).

Character	Mode	Turns	Content words	Word length (SD)	Zipf	Subjectivity	*S* within	*S* between
O-lan	First	87	848	5.94 (2.03)	4.98	0.434	0.943	0.945
	Third	82	1020	6.43 (2.12)	4.81	0.466		
Wang Lung	First	75	689	5.56 (1.90)	5.01	0.462	0.937	0.941
	Third	72	941	6.40 (2.12)	4.77	0.451		
the eldest son	First	62	466	5.81 (2.14)	5.02	0.494	0.932	0.942
	Third	56	694	6.47 (2.29)	4.85	0.556		
the maid	First	29	227	5.70 (2.19)	4.98	0.446		
	Third	24	234	6.50 (2.36)	4.81	0.440		

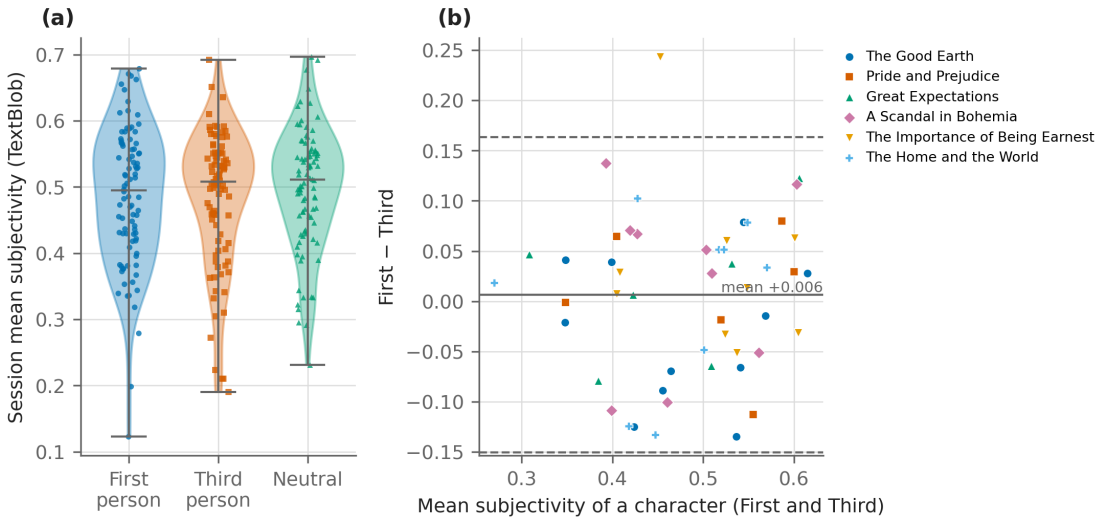


Figure 6. Subjectivity by narrative mode. (a) Session-level mean subjectivity in the three modes ($n = 90$ sessions each). (b) Bland–Altman plot of character-level subjectivity in first- versus third-person sessions; solid line: mean difference; dashed lines: 95% limits of agreement.

another. Figure 7a compares each character's cross-mode consistency with its consistency with the other characters of the scene. The two were practically identical (mean $S = 0.936$ within vs. 0.933 between characters; Wilcoxon $p = 0.347$, $n = 47$ characters), so the distributional measure reflects the shared style of the generator rather than character identity.

Speaker attribution is more sensitive. Within a mode, the content-word classifier identified the speaker with a balanced accuracy of 0.84 (chance 0.38). Across modes, accuracy dropped to 0.63, still above chance in 18 of 18 cells (Wilcoxon $p < 0.001$) but below the within-mode accuracy ($p < 0.001$; Fig. 7b); function words gave the same pattern at a lower level (0.69 within and 0.51 across modes). Character identity is therefore partially preserved: the agents keep their topics and concerns, while the narrative frame changes how they are expressed.

Turn-level subjectivity fluctuated around a stable level in all three modes, without a systematic trend or a separation between the modes (Fig. 8).

3.7 Ablation study: narrative control, memory and speaker selection

To isolate narrative control, actor memory and speaker selection, we compared ablated versions with the full framework on the same excerpts, modes and seeds (Llama 3.1 8B; Table 7, Fig. 9) and added a single-prompt baseline in which one LLM call writes the whole continuation.

Narrative control. Without a perspective instruction, the models rendered the turns as first-person speech, so the instruction is necessary to obtain third-person narration. Without the Director's compliance check and rewriting, 10.3% of the third-person turns violated the mode, and the first-minus-third difference in the first-person ratio fell from 0.898 to 0.833 ($p = 0.002$). The single-prompt baseline hardly separated the modes (difference 0.350), because many *third-person* turns were still direct speech. Turn-level control of narrative voice therefore requires the multi-agent design with explicit enforcement.

Actor memory. A memory of one turn ($k = 1$) did not change subjectivity ($p = 0.704$) or lexical diversity ($p = 0.907$), but it lowered compliance at the first attempt (0.899 vs. 0.957, $p = 0.003$): the recent turns also serve as in-context examples of the required form.

Speaker selection. With the addressee-aware policy, 90.8% of the turns came from the character addressed in the previous turn, compared with 34.0% under uniform random selection and 66.4% under a round-robin order (both $p < 0.001$); uniform selection also let a character speak twice in a row in 46.5% of the turns. No policy changed the narrative-mode effects (Table 7).

Judged quality. Both LLM judges rated the sessions with uniform speaker selection as the least coherent (Phi-4: 4.25 vs. 4.39 for the full framework; Qwen3.5: 3.14 vs. 3.61) and the single-prompt baseline as the most coherent (4.69 and 3.92). The multi-agent design is therefore not needed for fluency, but for turn-level control of voice, speaker and addressee.

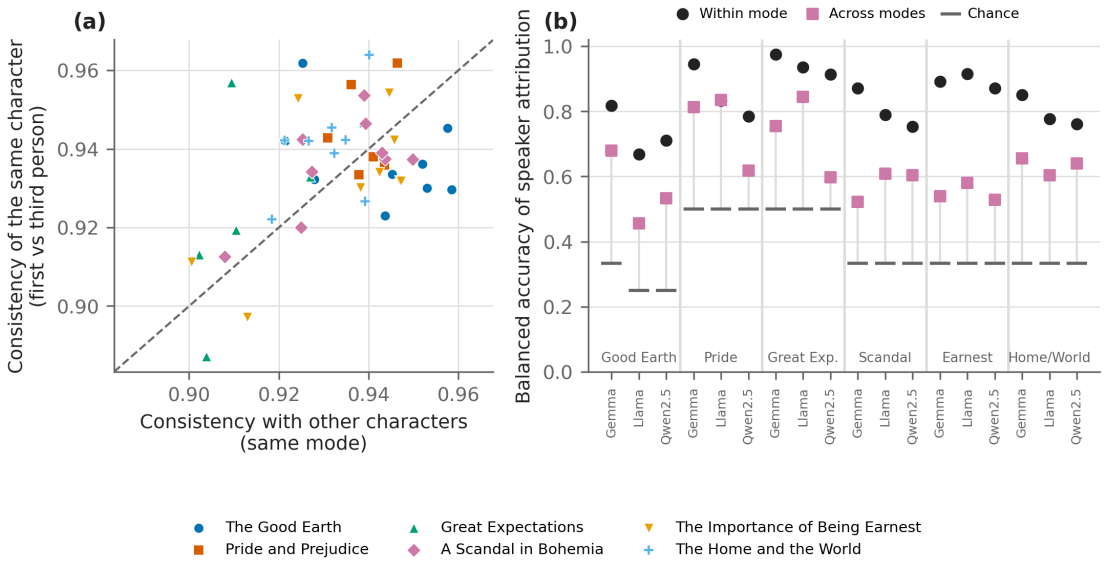


Figure 7. Character identity across narrative modes. (a) Composite distributional consistency *S* of each character between its first- and third-person turns (y) against its mean consistency with the other characters of the scene in the same mode (x); points above the diagonal would indicate preserved character-specific traits. (b) Balanced accuracy of speaker attribution from content words within a mode (leave-one-seed-out) and across modes (trained on one mode, tested on the other) for every excerpt–model cell; horizontal bars mark chance.

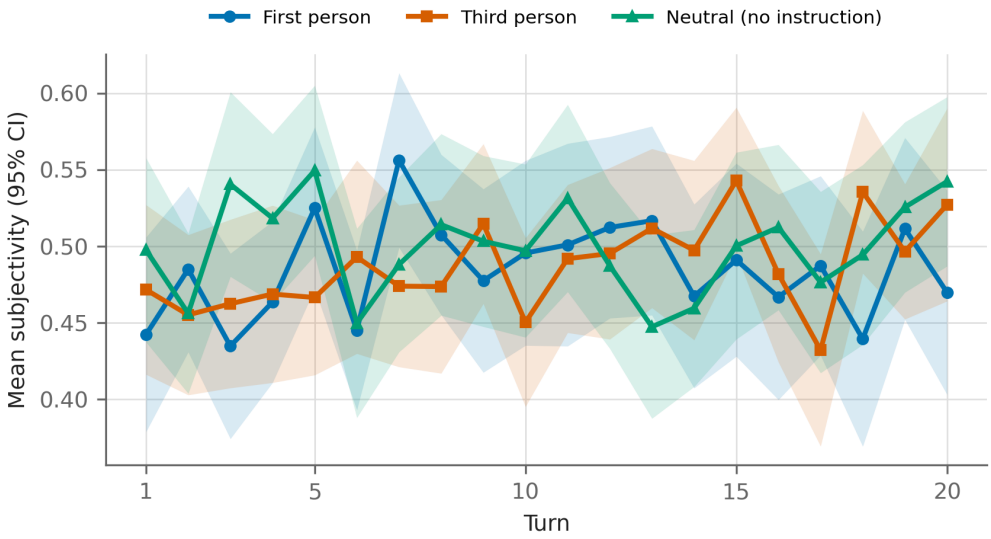


Figure 8. Turn-level subjectivity by narrative mode (mean over all excerpts, models and seeds with 95% confidence bands).

Table 7. Ablation results (Llama 3.1 8B; 6 excerpts × 2 modes × 3 seeds per condition), compared with the full framework on the same excerpts, modes and seeds. Responsiveness: share of turns by the character addressed in the previous turn; turn balance: normalized entropy of the speaker distribution; self-succession: share of consecutive turns by the same speaker; repetition: 1 − share of distinct word trigrams; compliant (1st) and residual viol.: share of turns that followed the mode at the first attempt and after enforcement. Gap: mean first − third difference of subjectivity and of the first-person ratio within a block. *Paired Wilcoxon test against the full framework, $p < 0.05/6$.

Condition	Responsiveness	Turn balance	Self-succession	Repetition	Compliant (1st)	Residual viol.	Gap subj.	Gap FP ratio
Full framework	0.908	0.955	0.007	0.150	0.957	0.000	−0.015	0.898
−Memory ($k = 1$)	0.938	0.935	0.006	0.140	0.899*	0.010	0.009	0.913
−Speaker selection (uniform)	0.340*	0.974	0.465*	0.131	0.944	0.001	0.015	0.912
−Speaker selection (round robin)	0.664*	0.999*	0.000	0.148	0.968	0.001	0.032	0.895
−Mode enforcement	0.904	0.956	0.009	0.139*	0.949	0.051*	−0.029	0.833
Single-prompt baseline	0.808	0.968	0.004	0.116	–	–	0.049	0.350

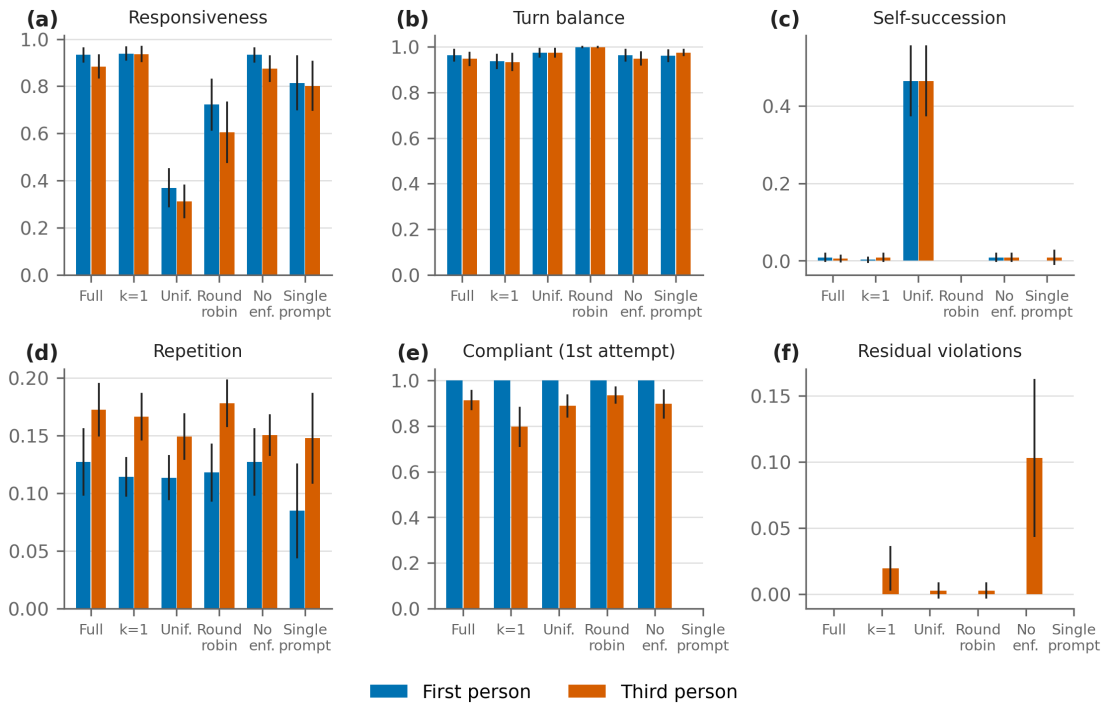


Figure 9. Ablation of memory, speaker selection and mode enforcement, and the single-prompt baseline (Llama 3.1 8B, 6 excerpts, 3 seeds; blue: first-person sessions, orange: third-person sessions; error bars: 95% confidence intervals). (a) Addressee responsiveness; (b) turn-taking balance; (c) self-succession; (d) repetition; (e) mode compliance at the first attempt; (f) residual mode violations after enforcement. Full: full framework; $k = 1$: memory of one turn; Unif.: uniform speaker selection; Round robin: fixed speaker order; No enf.: without Director enforcement; Single prompt: single-prompt baseline.

3.8 LLM-judge evaluation

Table 8 and Fig. 10 summarize the judges' ratings. The judges agreed only weakly (Krippendorff's $\alpha = 0.41$ for subjectivity, 0.23 for coherence and 0.01 for emotional intensity; negative for character consistency and scene fidelity, which Phi-4 rated near the maximum), so their ratings are interpreted separately.

Qwen3.5 classified the narrative form of 100.0% of the first-person and 100.0% of the third-person sessions as instructed; Phi-4 labelled many third-person sessions as "mixed" (46.7% correct). Phi-4 found no difference between the modes in perceived subjectivity ($p = 0.618$) or emotional intensity ($p = 0.209$). Qwen3.5 rated third-person sessions lower on every criterion, including character consistency and scene fidelity (all $p < 0.001$), which should not depend on the narrative form; this suggests a general preference for direct-speech dialogue rather than a specific perception of lower subjectivity. The judges' subjectivity ratings correlated only weakly with TextBlob subjectivity (Spearman $\rho = 0.13$ and 0.12 at the session level), so perceived subjectivity is measure-dependent.

In the masked speaker-identification task, both judges identified the speaker far above chance in both modes (Phi-4: 0.89 for first-person and 0.88 for third-person turns; Qwen3.5: 0.91 and 0.89; chance 0.38; binomial tests $p < 0.001$). With the role cards available, the turns thus remain tied to their characters' situations and concerns in both modes.

Table 8. LLM-judge evaluation of the main-experiment sessions (1–5 scales; $n = 90$ sessions per mode and judge). Form correct: share of sessions whose dominant narrative form was classified as the instructed one. Speaker ID: accuracy of identifying the speaker of a turn with all names masked (chance level in parentheses).

Judge	Mode	Coherence	Char. consistency	Scene fidelity	Subjectivity	Emot. intensity	Form correct	Speaker ID (chance)
Phi-4 14B	First	4.73	5.00	4.89	3.62	3.86	0.69	0.89 (0.38)
Phi-4 14B	Neutral	4.71	5.00	4.89	3.60	3.76	–	0.92 (0.38)
Phi-4 14B	Third	4.62	4.99	4.83	3.59	3.78	0.47	0.88 (0.37)
Qwen3.5 9B	First	4.53	3.64	2.89	4.16	3.47	1.00	0.91 (0.38)
Qwen3.5 9B	Neutral	4.54	3.64	2.84	4.20	3.52	–	0.93 (0.38)
Qwen3.5 9B	Third	3.74	2.43	1.96	2.93	2.48	1.00	0.89 (0.37)

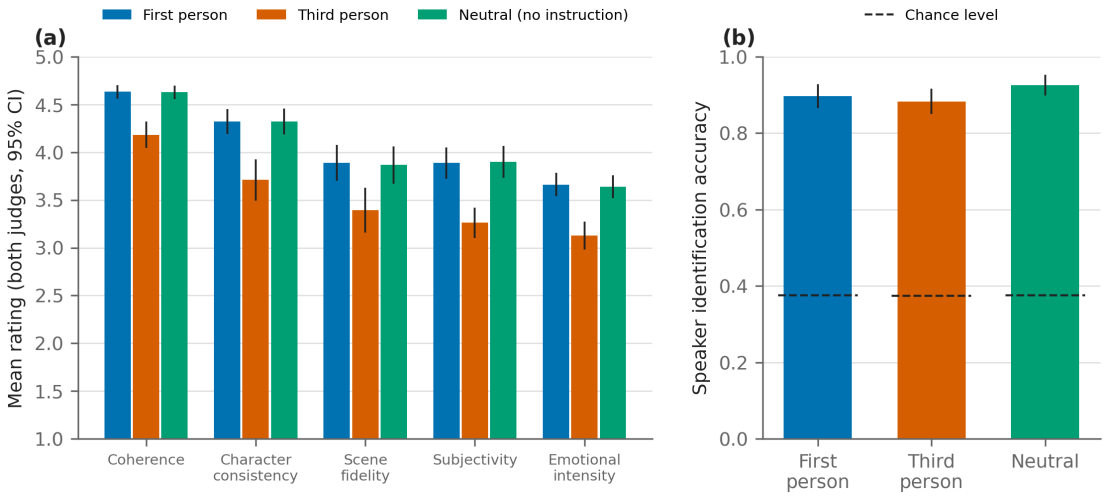


Figure 10. LLM-judge evaluation. (a) Mean ratings of the main-experiment sessions by narrative mode (both judges pooled; error bars: 95% confidence intervals). (b) Accuracy of masked speaker identification by narrative mode (both judges pooled); dashed lines mark chance.

3.9 Improving the preservation of character identity

The standard third-person mode loses part of the characters' identity and shifts their vocabulary towards the narrator's register. We therefore tested three ways of keeping the characters' voice in third-person narration. All three change only the narrative-mode component and are compared with the standard

framework on the same excerpts, generators and seeds (Table 9; paired comparisons in Supplemental Information, S11).

Voice exemplars. Adding up to three of each character's lines from the excerpt to the role card (Appendix A.7) did not make the characters more distinguishable (cross-mode attribution 0.60 vs. 0.60 without exemplars) and lowered compliance at the first attempt (0.957 vs. 0.978), because the exemplars, which are direct speech, pulled third-person turns back towards direct speech.

Voice-preserving instruction. Asking the narrator to keep the character's own words, idioms and tone (free indirect style) raised cross-mode attribution from 0.63 to 0.67 ($p = 0.074$) and reduced the register shift in content-word length from 0.46 to 0.38 characters. However, only 67.7% of the turns complied with the third-person form at the first attempt, 4.0% still contained direct speech after enforcement, and both judges rated the sessions as less coherent.

Speak-then-narrate. The third variant separates what a character says from how the narrator reports it. The Actor receives the first-person instruction and speaks in the character's own voice; the Director then turns the speech into third-person narration in free indirect style, changing only pronouns, tense and at most a short reporting phrase (Appendix A.7). Because the Actor prompt is identical to that of the first-person condition, this variant isolates the contribution of the narration itself. Mode control remained reliable (0.6% residual violations), and cross-mode attribution rose from 0.60 to 0.73 (higher in 17 of 18 cells, $p < 0.001$). The register shift largely disappeared: narrated turns no longer differed reliably from first-person speech in content-word length ($p_{\text{Holm}} = 0.364$), word frequency ($p_{\text{Holm}} = 1.000$) or emotional intensity ($p_{\text{Holm}} = 1.000$). The LLM judges, however, identified the speakers of narrated turns less accurately than those of standard third-person turns (Phi-4: 0.78 vs. 0.88; Qwen3.5: 0.81 vs. 0.89) and rated the sessions as less coherent (Phi-4: 4.30 vs. 4.61; Qwen3.5: 2.22 vs. 3.67), possibly because turns that keep the character's wording without explanatory framing read as less natural.

Summary. What matters for character identity is who formulates the third-person turn. When the narrator paraphrases it, the vocabulary shifts towards the narrator's register; when the narration only transposes the character's words, the lexical identity of the characters is largely preserved, at the price of lower judged coherence. The more formal register of third-person narration in our simulations is thus produced by narratorial paraphrase rather than by the grammatical person (content-word length, standard vs. narrated third person: 7.04 vs. 6.73, $p_{\text{Holm}} < 0.001$). We keep the standard mode as the main condition, because it corresponds to conventional narratorial report, and provide speak-then-narrate as an option for studies that need to hold the characters' wording constant.

Table 9. Attempts to preserve character identity in third-person narration. Each cell gives the value of the variant and, in parentheses, the value of the standard framework on the same excerpts, models and seeds. Compliance and residual violations refer to third-person turns (voice exemplars: both modes). Cross-mode attribution: balanced accuracy of a content-word classifier trained on one mode and tested on the other (first person vs. the variant's third-person mode). Speaker ID: accuracy of the LLM judges on masked third-person turns (seeds 0–2). |FP−TP|: mean absolute first-minus-third difference per excerpt, model and seed (content-word length in characters; emotional intensity).

Variant	Sessions	Compliant (1st)	Residual viol.	Cross-mode attribution	Speaker ID Phi-4	Speaker ID Qwen3.5	|FP−TP| word length	|FP−TP| intensity
Voice exemplars (both modes)	108	0.957 (0.978)	0.002 (0.000)	0.60 (0.60)	0.86 (0.88)	0.86 (0.89)	–	–
Voice-preserving instruction	90	0.677 (0.958)	0.040 (0.000)	0.67 (0.63)	0.87 (0.88)	0.89 (0.89)	0.38 (0.46)	0.054 (0.063)
Speak-then-narrate	54	0.894 (0.958)	0.006 (0.000)	0.73 (0.60)	0.78 (0.88)	0.81 (0.89)	0.26 (0.47)	0.044 (0.055)

3.10 Expert evaluation

To check the automatic measures and the LLM judges against human readers, three raters with master's degrees in English literature (two) or English drama (one) evaluated a blinded subset of 12 sessions: four excerpts (*Pride and Prejudice*, *Great Expectations*, "A Scandal in Bohemia" and *The Home and the World*, each with one generator), each in the first-person, standard third-person and speak-then-narrate conditions (seed 0, first eight turns), together with 36 masked turns from those rated by the LLM judges. The raters did not know the generator or the condition. They rated every session on the judges' 1–5 scales for coherence, character consistency, subjectivity and emotional intensity, classified its narrative form, and identified the speaker of every masked turn. The instructions were in Chinese and the dialogues in English; the raters were asked to work independently and without AI tools. Their responses were entered into the rating sheets by one of the authors, who was not blind to the conditions; the item-level ratings are

released with the data (Supplemental Information, S12).

The experts agreed moderately to substantially (Krippendorff's α from 0.55 for emotional intensity to 0.79 for coherence), more than the LLM judges ($\alpha \leq 0.41$), although on far fewer sessions. They identified the speakers of masked turns far above chance in all three conditions (first person: 1.00, standard third person: 0.86, speak-then-narrate: 0.97; chance 0.42; binomial tests of the majority vote $p = 0.004$ or smaller) and more accurately than the judges on the same turns (Phi-4: 0.75; Qwen3.5: 0.75). Like the judges, they rated the speak-then-narrate sessions as the least coherent (2.33, vs. 3.25 for first-person and 2.83 for standard third-person sessions), whereas perceived subjectivity differed little between the conditions (4.08, 3.83 and 4.17). Their session ratings correlated moderately with those of the judges (Spearman ρ from 0.45 to 0.69) and with the automatic measures of subjectivity ($\rho = 0.68$) and emotional intensity ($\rho = 0.43$). Because each condition mean rests on four excerpts and the responses were transcribed by an author, we report these results descriptively, as a plausibility check rather than a test of the hypotheses.

4 DISCUSSION

4.1 Interpretation of Findings

Our experiments support some of the hypotheses that motivated this study and not others. First, the framework controlled narrative voice reliably across six excerpts and three generators. Second, the clearest consequence of switching from first-person speech to third-person narration was a more formal lexical register, consistent across texts and models and opposite to what pronoun substitution produces. Third, third-person narration was affectively more neutral; the comparison with the no-instruction control shows that narration dampens affect rather than first-person speech amplifying it, and the effects were small to moderate. Fourth, lexicon-based subjectivity was unaffected, the LLM judges disagreed on perceived subjectivity, and the experts perceived little difference between the modes. In our simulations, these results therefore support the hypothesis that third-person narration promotes a more neutral, descriptive register, support stronger emotional polarity in first-person speech only in this relative sense, and do not support the hypothesis that first-person narration increases subjectivity, at least as measured here.

At the character level, the Jensen–Shannon measure did not distinguish characters, so it cannot show that traits are preserved, and a small divergence would in any case characterize only the simulations produced here, not the approach in general. Speaker attribution, the judges and the experts indicate that character identity is partially preserved across modes. The improvement experiments locate the loss in the narrator's paraphrase: when the narration only transposes the character's own speech, cross-mode attribution rises and the register shift largely disappears.

4.2 Prompt design and the interpretation of perspective effects

A central concern for any prompt-based manipulation is that differences between modes may be direct consequences of the prompts. Our analysis separates three kinds of effects. (i) Effects that the instruction states explicitly, such as the pronoun profile, are manipulation checks. (ii) Effects of rendering a turn as reported speech are estimated by the rule-based transposition, which accounts for most of the difference in utterance length and for part of the DistilBERT difference, but not for the differences in content-word length, word frequency, intensity and polarity. (iii) Only effects beyond the transposition indicate how the generators realize the instruction to narrate; even these describe LLMs instructed to write in a given mode and may depend on the wording of the instructions. The source texts provide a reference point: in the five prose excerpts, narration had higher TextBlob subjectivity (0.379) than quoted speech (0.267) and longer content words (6.14 vs. 5.49). Lexicon-based subjectivity thus mainly reflects evaluative and descriptive adjectives, which narrators use at least as much as characters, rather than the experiential subjectivity that narratologists associate with first-person narration; the absence of a subjectivity effect is therefore partly a limitation of the measure.

4.3 Strengths and Limitations

The framework integrates character modeling, narrative-mode control and quantitative analysis into a replicable pipeline: prompts are fixed, random choices are seeded, model responses are cached, and the code, dialogues and analysis scripts are released. The paired design keeps role cards, opening and seed constant across modes, and the conclusions rest on six texts, three models and five seeds with paired tests and correction for multiple comparisons.

Several limitations remain. First, the results describe three 4-bit quantized open-weight models of 8–14B parameters; larger or proprietary models may realize the modes differently. Second, the excerpts are short single scenes in English (one in translation), so the findings may not extend to long-form narrative, other languages or other scenes. Third, the automatic measures were developed for other domains, and their validity for literary dialogue is limited, as the comparison with the source texts shows. Fourth, the LLM judges are not a substitute for human evaluation: one belongs to the same model family as a generator, they agreed only weakly (Krippendorff's $\alpha \leq 0.41$), and one preferred direct-speech dialogue in general. The expert evaluation was small (3 raters and 12 sessions), its responses were transcribed by an author who was not blind to the conditions, and it did not include lay readers. Fifth, the cast lists were verified manually, and the Director's extraction still misattributes some quotations. Finally, the modes manipulate grammatical person and narratorial mediation, not focalization.

Narratological implications. Within these limits, the simulations suggest hypotheses that can be tested with human-authored texts and readers: narratorial mediation is expressed through a more formal vocabulary and a moderate dampening of affect rather than through a change of lexicon-based subjectivity; changing the narrative voice preserves what characters talk about more than how they say it, unless the narrator stays close to the characters' own words, as in free indirect style; and first-person direct speech is the default voice of role-playing LLMs, which should be considered when such models are used to simulate narrative.

5 CONCLUSION AND FUTURE WORK

We have presented an LLM-driven multi-agent simulation framework for studying narrative voice shifts under controlled conditions, in which a Director Agent enforces the narrative mode and Actor Agents with a sliding-window memory generate the dialogue. In experiments with six excerpts by six authors, three open-weight models and five seeds, the framework reliably controlled narrative voice. In our simulations, third-person narration was associated with a more formal vocabulary and lower emotional intensity than first-person speech, whereas lexicon-based subjectivity did not change, and character identity was partially preserved across modes. A speak-then-narrate variant preserved character identity much better and removed most of the register shift, which indicates that the shift arises from narratorial paraphrase. These results support some of the hypotheses that motivated this study and not others, and they show why single-run results from one text and one model should not be generalized.

In our future work, we will conduct a larger evaluation with literary experts and lay readers, test larger models, longer texts and other languages, and develop measures of narratorial subjectivity and focalization that suit literary text better than general-purpose sentiment lexicons.

ACKNOWLEDGMENTS

The authors would like to express their sincere gratitude to Dr. Karry Zhang for his invaluable support and guidance throughout the development of this work.

REFERENCES

Abdin, M., Aneja, J., Behl, H., Bubeck, S., Eldan, R., Gunasekar, S., Harrison, M., Hewett, R. J., Javaheripi, M., Kauffmann, P., Lee, J. R., Lee, Y. T., Li, Y., Liu, W., Mendes, C. C. T., Nguyen, A., Price, E., de Rosa, G., Saarikivi, O., Salim, A., Shah, S., Wang, X., Ward, R., Wu, Y., Yu, D., Zhang, C., and Zhang, Y. (2024). Phi-4 technical report. *arXiv preprint arXiv:2412.08905*.

Austen, J. (1813). *Pride and Prejudice*. T. Egerton, London. Project Gutenberg eBook #1342.

Bal, M. (2009). *Narratology: Introduction to the Theory of Narrative*. University of Toronto Press, 3rd edition.

Baroni, R. and Paschoud, A. (2021). Introduction: Time and narrative, the missing link between the "narrative turn" and postclassical narratology?

Bland, J. M. and Altman, D. G. (1986). Statistical methods for assessing agreement between two methods of clinical measurement. *The Lancet*, 327(8476):307–310.

Brooke, J., Hammond, A., and Hirst, G. (2017). Using models of lexical style to quantify free indirect discourse in modernist fiction. *Digital Scholarship in the Humanities*, 32(2):234–250.

Brunyé, T. T., Ditman, T., Mahoney, C. R., Augustyn, J. S., and Taylor, H. A. (2009). When you and I share perspectives: Pronouns modulate perspective taking during narrative comprehension. *Psychological Science*, 20(1):27–32.

Buck, P. S. (1931). *The Good Earth*. John Day, New York.

Cai, J., Li, J., Zhang, M., Li, M., Wang, C.-S., and Tei, K. (2024). Language evolution for evading social media regulation via llm-based multi-agent simulation. In *IEEE Congr. Evol. Comput. (CEC)*, pages 1–10. IEEE.

Cavazza, M., Charles, F., and Mead, S. J. (2002). Character-based interactive storytelling. *IEEE Intell. Syst.*, 17(4):17–24.

Chiang, C.-H. and Lee, H.-y. (2023). Can large language models be an alternative to human evaluations? In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 15607–15631. Association for Computational Linguistics.

Covington, M. A. and McFall, J. D. (2010). Cutting the Gordian knot: The moving-average type–token ratio (MATTR). *Journal of Quantitative Linguistics*, 17(2):94–100.

De Zarzà, I., De Curtò, J., Roig, G., Manzoni, P., and Calafate, C. T. (2023). Emergent cooperation and strategy adaptation in multi-agent systems: An extended coevolutionary theory with llms. *Electronics*, 12(12):2722.

Dickens, C. (1861). *Great Expectations*. Chapman and Hall, London. Project Gutenberg eBook #1400.

Doyle, A. C. (1892). *The Adventures of Sherlock Holmes*. George Newnes, London. "A Scandal in Bohemia" first published in The Strand Magazine, 1891; Project Gutenberg eBook #1661.

Efron, B. and Tibshirani, R. J. (1993). *An Introduction to the Bootstrap*. Chapman & Hall, New York.

Eisenberg, J. and Finlayson, M. (2016). Automatic identification of narrative diegesis and point of view. In *Proceedings of the 2nd Workshop on Computing News Storylines (CNS 2016)*, pages 36–46, Austin, TX. Association for Computational Linguistics.

Fludernik, M. (2009). *An Introduction to Narratology*. Routledge, London.

Gemma Team (2025). Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*.

Genette, G. (1980). *Narrative Discourse: An Essay in Method*. Cornell University Press, Ithaca, NY. Translated by Jane E. Lewin.

Gius, E. and Vauth, M. (2022). Towards an event based plot model: A computational narratology approach. *J. Comput. Lit. Stud.*, 1(1).

Guo, T., Chen, X., Wang, Y., Chang, R., Pei, S., Chawla, N. V., Wiest, O., and Zhang, X. (2024). Large language model based multi-agents: A survey of progress and challenges. In *Proc. Int. Joint Conf. Artif. Intell.*, pages 8048–8057.

Hartung, F., Burke, M., Hagoort, P., and Willems, R. M. (2016). Taking perspective: Personal pronouns affect experiential aspects of literary reading. *PLoS ONE*, 11(5):e0154732.

Hatavara, M., Sandberg, K., Andrushchenko, M., Hälikkö, S., Nummenmaa, J., Nummenmaa, T., Peltonen, J., and Hyvärinen, M. (2024). Computational recognition of narratives: Applying narratological definitions to the analysis of political language use. *Narrative Inq.*, 34(2):335–363.

Hayes, A. F. and Krippendorff, K. (2007). Answering the call for a standard reliability measure for coding data. *Communication Methods and Measures*, 1(1):77–89.

He, X., Lashkari, A. H., Vombatkere, N., and Sharma, D. P. (2024). Authorship attribution methods, challenges, and future research directions: A comprehensive survey. *Information*, 15(3):131.

Herrmann, J. B., Jacobs, A. M., and Piper, A. (2021). Computational stylistics. *Handb. Empir. Lit. Stud.*, pages 451–486.

Holm, S. (1979). A simple sequentially rejective multiple test procedure. *Scandinavian Journal of Statistics*, 6(2):65–70.

Honnibal, M., Montani, I., Van Landeghem, S., and Boyd, A. (2020). spaCy: Industrial-strength natural language processing in Python.

Hutto, C. J. and Gilbert, E. (2014). VADER: A parsimonious rule-based model for sentiment analysis of social media text. In *Proceedings of the Eighth International AAAI Conference on Weblogs and Social Media*, pages 216–225.

Khafaga, A. F. and Shaalan, I. E.-N. A. W. (2022). A computational approach to decode the pragma-stylistic meanings in narrative discourse. *Int. J. Adv. Comput. Sci. Appl.*, 13(2).

Li, X., Wang, S., Zeng, S., Wu, Y., and Yang, Y. (2024). A survey on llm-based multi-agent systems: Workflow, infrastructure, and challenges. *Vicinagearth*, 1(1):9.

Lin, J. (1991). Divergence measures based on the Shannon entropy. *IEEE Transactions on Information Theory*, 37(1):145–151.

Llama Team, AI @ Meta (2024). The Llama 3 herd of models. *arXiv preprint arXiv:2407.21783*.

Loria, S. (2020). TextBlob: Simplified text processing. https://textblob.readthedocs.io. Python library.

Magerko, B., Manzoul, W., Riedl, M., Baumer, A., Fuller, D., Luther, K., and Pearce, C. (2009). An empirical study of cognition and theatrical improvisation. In *Proc. 7th ACM Conf. Creativity and Cognition*, pages 117–126.

Mateas, M. and Stern, A. (2003). Façade: An experiment in building a fully-realized interactive drama. In *Proceedings of the Game Developers Conference, Game Design Track*, San Jose, CA.

McCarthy, P. M. and Jarvis, S. (2010). MTLD, vocd-D, and HD-D: A validation study of sophisticated approaches to lexical diversity assessment. *Behavior Research Methods*, 42(2):381–392.

Niederhoff, B. (2013). Focalization. In Hühn, P., Meister, J. C., Pier, J., and Schmid, W., editors, *The Living Handbook of Narratology*. Hamburg University, Hamburg.

Ollama (2026). Ollama: Run large language models locally (version 0.40.3). https://ollama.com. Accessed 13 October 2026.

Onega, S. and Landa, J. A. G. (1996). *Narratology: An Introduction*. Addison Wesley Longman.

Park, J. S., O'Brien, J., Cai, C. J., Morris, M. R., Liang, P., and Bernstein, M. S. (2023). Generative agents: Interactive simulacra of human behavior. In *Proceedings of the 36th Annual ACM Symposium on User Interface Software and Technology (UIST '23)*, pages 1–22. ACM.

Piper, A., So, R. J., and Bamman, D. (2021). Narrative theory for computational narrative understanding. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 298–311. Association for Computational Linguistics.

Porteous, J. and Cavazza, M. (2009). Controlling narrative generation with planning trajectories: The role of constraints. In *Interactive Storytelling*, volume 5915 of *Lect. Notes Comput. Sci.*, pages 234–245, Berlin, Heidelberg. Springer.

Qwen Team (2024). Qwen2.5 technical report. *arXiv preprint arXiv:2412.15115*.

Qwen Team (2026). Qwen3.5 model family. Ollama model library, https://ollama.com/library/qwen3.5. Accessed 13 October 2026.

Rahimi, H., Cattoni, J., Beghili, M., Abrini, M., Khoramshahi, M., Pino, M., and Chetouani, M. (2025). Reasoning llms for user-aware multimodal conversational agents. In *2025 34th IEEE International Conference on Robot and Human Interactive Communication (RO-MAN)*. IEEE.

Riedl, M. O. and Young, R. M. (2010). Narrative planning: Balancing plot and character. *J. Artif. Intell. Res.*, 39:217–268.

Sanh, V., Debut, L., Chaumond, J., and Wolf, T. (2019). DistilBERT, a distilled version of BERT: Smaller, faster, cheaper and lighter. *arXiv preprint arXiv:1910.01108*.

Shanahan, M., McDonell, K., and Reynolds, L. (2023). Role play with large language models. *Nature*, 623(7987):493–498.

Shao, Y., Li, L., Dai, J., and Qiu, X. (2023). Character-LLM: A trainable agent for role-playing. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 13153–13187. Association for Computational Linguistics.

Si, M., Marsella, S. C., and Pynadath, D. V. (2005). Thespian: Using multi-agent fitting to craft interactive drama. In *Proc. 4th Int. Joint Conf. Auton. Agents Multiagent Syst.*, pages 21–28.

Speer, R. (2022). rspeer/wordfreq: v3.0.

Sun, C., Huang, S., and Pompili, D. (2025). Llm-based multi-agent decision-making: Challenges and future directions. *IEEE Robot. Autom. Lett.*, PP(99):1–8.

Swartjes, I. and Theune, M. (2006). A fabula model for emergent narrative. In *Technologies for Interactive Digital Storytelling and Entertainment*, Lect. Notes Comput. Sci., pages 49–60. Springer.

Tagore, R. (1919). *The Home and the World*. Macmillan, London. Translated by Surendranath Tagore; Project Gutenberg eBook #7166.

Vishnubhotla, K., Hammond, A., and Hirst, G. (2022). The project dialogism novel corpus: A dataset for quotation attribution in literary texts. In *Proceedings of the Thirteenth Language Resources and Evaluation Conference*, pages 5838–5848, Marseille, France. European Language Resources Association.

Wang, L., Lian, J., Huang, Y., Dai, Y., Li, H., Chen, X., Xie, X., and Wen, J.-R. (2024). Characterbox:

Evaluating the role-playing capabilities of llms in text-based virtual worlds. *arXiv*.
Wang, S. (2025). Enhancing in-context learning of large language models for knowledge graph reasoning via rule-and-reinforce selected triples. *Appl. Sci.*, 15(3):1088.
Wilcoxon, F. (1945). Individual comparisons by ranking methods. *Biometrics Bulletin*, 1(6):80–83.
Wilde, O. (1899). *The Importance of Being Earnest: A Trivial Comedy for Serious People*. Leonard Smithers, London. First performed 1895; Project Gutenberg eBook #844.
Yu, Y., Yu, R., Wei, H., Zhang, Z., and Qian, Q. (2025). Beyond dialogue: A profile-dialogue alignment framework towards general role-playing language model. In *Proc. 63rd Annu. Meet. Assoc. Comput. Linguist. (Vol. 1: Long Pap.)*.
Zhai, Y., Yang, T., Xu, K., Feng, D., Yang, C., Ding, B., and Wang, H. (2025). Enhancing decision-making for llm agents via step-level q-value models. In *Proc. AAAI Conf. Artif. Intell.*, volume 39, pages 27161–27169.
Zhang, P., An, S., Qiao, L., Yu, Y., Chen, J., Wang, J., Yin, D., Sun, X., and Zhang, K. (2025a). Roleplot: A systematic framework for evaluating and enhancing the plot-progression capabilities of role-playing agents. In *Proc. 63rd Annu. Meet. Assoc. Comput. Linguist. (Vol. 1: Long Pap.)*, pages 12337–12354.
Zhang, T., Liu, Z., Xin, Y., and Jiao, Y. (2025b). Mooseagent: A llm based multi-agent framework for automating moose simulation. *arXiv*.
Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E. P., Zhang, H., Gonzalez, J. E., and Stoica, I. (2023). Judging LLM-as-a-judge with MT-Bench and chatbot arena. In *Advances in Neural Information Processing Systems 36 (Datasets and Benchmarks Track)*.

APPENDIX: PROMPT TEMPLATES

All prompts are fixed templates; text in angle brackets is filled in by the framework. JSON output is additionally enforced with the JSON mode of the Ollama API. The complete source code is available in the project repository.

A.1 Dialogue extraction (Director, temperature 0).

```
Extract the direct-speech dialogue from the following literary excerpt, in order of appearance.
Rules:
1. Include only spoken utterances (no narration, inner thoughts or action descriptions).
2. If an utterance is interrupted by narration, merge the parts into one utterance.
3. "speaker": use the character's name if it is known from the excerpt; otherwise the first
     descriptive reference (e.g. "the old lady"). Use exactly the same label for the same person
     throughout. Never invent labels such as "Speaker A".
4. "speak_to": always a non-empty array of addressee labels (use the same labels as for speakers);
      infer the addressee from context if it is not explicit.
Return only JSON of the form {"dialogues": [{"speaker": "...", "speak_to": ["..."], "text":
     "..."}]}.
Excerpt:
<excerpt>
```

A.2 Role cards (Director, temperature 0).

```
You are the Director of a dialogue simulation. Based on the literary excerpt and its extracted
     dialogue, build a role card for every character who speaks.
Return only JSON of the form:
{"characters": [{"name": "...", "gender": "female|male|unknown", "personality": "...", "
     emotional_baseline": 0.0, "speech_style": "...", "description": "...", "relationships": {"
     Other name": "relation"}}],
 "context": "two or three sentences summarising the situation of the scene"}
"emotional_baseline" is a number in [-1, 1]. Use exactly these speaker labels as names: <cast
     names>.
Excerpt: <excerpt>
Extracted dialogue: <dialogue>
```

A.3 Narrative-mode instructions (inserted into the Actor prompt).

```
FIRST PERSON:  NARRATIVE MODE: FIRST PERSON. Write <name>'s turn as direct speech in <name>'s own
     voice. <name> refers to himself or herself with first-person pronouns (I, me, my). Do not add
     a narrator, quotation marks, or descriptions of <name> in the third person.
THIRD PERSON:  NARRATIVE MODE: THIRD PERSON. Do not write direct speech. Write <name>'s turn as
     narration by an external narrator, as in a novel: report what <name> says (and, if relevant,
     does or feels), naming <name> and using a reporting verb (told, asked, insisted, ...). Never
     use I, me, my, we, our, you or your, and no quotation marks. Example. Direct speech: John said
      to Mary, "I'm tired and don't want to go to the party." Third-person narration: John told
     Mary that he was tired and had no intention of going to the party.
NEUTRAL:       (no instruction)
```

A.4 Actor Agent (temperature 0.7, seeded).

```
You are role-playing the literary character <name> in a multi-character scene.
Role card:
- Name: <name>   - Personality: <personality>   - Speech style: <speech style>
- Background: <description>   - Relationships: <relationships>
- Current emotion (-1 negative .. 1 positive): <emotional baseline>
Scene: <context>
Other characters present: <other names>
Recent dialogue (oldest first):
<last k turns, one per line: "speaker (to addressees): text">
Instructions: Stay in character. Decide what <name> says next, consistent with the role card and
     the recent dialogue. Advance the scene; do not repeat earlier turns. Keep the turn to one to
     three sentences.
<narrative-mode instruction>
Return only JSON: {"intend": "speak" or "silent", "content": "<mode-specific description>", "
     speak_to": ["addressee name(s)"]}
```

When a turn violates the mode, the line “IMPORTANT: a previous answer did not follow the narrative mode (¡reason¿). Follow the NARRATIVE MODE exactly.” is appended to the mode instruction for the second attempt. If the second attempt also fails, the Director rewrites the turn with the prompt: “You are the Director of a dialogue simulation and you enforce the narrative mode. Rewrite the following turn of ¡name¿ (addressed to ¡addressees¿) so that it follows the mode below. Keep its meaning, tone and length; change only the narrative form. ¡mode instruction¿ Turn: ¡text¿”.

A.5 Single-prompt baseline (temperature 0.7, seeded).

```
Write the continuation of the following literary scene as a dialogue of exactly <n> turns among
     these characters: <names>.
Role cards: <one line per character>
Scene: <context>
Opening of the scene: <last five opening turns>
<mode text: first person = "Every turn must be direct speech in the speaker's own voice using
     first-person pronouns (I, me, my); no narrator and no quotation marks." / third person = "
     Every turn must be third-person narration by an external narrator reporting what the speaker
     says (by name or he/she, never 'I'), with no direct speech or quotation marks.">
Each turn is one to three sentences. Return only JSON of the form {"turns": [{"speaker": "...", "
     speak_to": ["..."], "text": "..."}]} with <n> items.
```

A.6 LLM judges (temperature 0).

```
SESSION TASK: You are an expert in literary studies evaluating a computer-generated continuation
     of a scene from "<title>" by <author>. Characters: <role cards>. Scene context: <context>.
     Continuation (<n> turns): <numbered turns>.
Rate the continuation on each criterion from 1 (very poor / very low) to 5 (excellent / very high):
      coherence (the turns follow sensibly from one another and form a meaningful exchange);
     character_consistency (each character's turns fit the character description and stay
     consistent); scene_fidelity (the continuation is plausible for this scene of the work:
     situation, relationships, period); subjectivity (how strongly the turns foreground personal
     opinions, evaluations and feelings; 1 = detached and factual, 5 = highly personal);
     emotional_intensity (1 = flat, 5 = intense). Also classify the dominant narrative form as "
     first_person_direct_speech", "third_person_narration" or "mixed". Return only JSON.

SPEAKER TASK: ... One turn of a generated dialogue is shown below. All character names were
     replaced by [NAME] and gendered pronouns by neutral ones, so you must decide from the content
     and style alone. Turn: <masked turn>
1) Which character most probably produced this turn? Answer with one name from this list: <names>.
2) Rate the subjectivity of the turn from 1 (detached, factual) to 5 (highly personal, evaluative,
      emotional). Return only JSON: {"speaker": "...", "subjectivity": 1-5}
```

A.7 Extension and improvement experiments.

```
VOICE EXEMPLARS (added to the Actor's role card): Voice examples (lines that <name> speaks in the
     source text; imitate the voice and diction, not the content):
- "<line 1>"  - "<line 2>"  - "<line 3>"
VOICE-PRESERVING THIRD PERSON (mode instruction of improvement 1): NARRATIVE MODE: THIRD PERSON,
     VOICE-PRESERVING. Do not write direct speech. Write <name>'s turn as narration by an external
     narrator, referring to <name> by name or he/she, never as "I", and use no quotation marks or
     the pronouns I, me, my, we, our, you or your. Keep <name>'s own words, idioms, tone and
     characteristic expressions as far as possible (free indirect style): report the turn in <name
     >'s voice instead of paraphrasing it in the narrator's formal language. Example. Direct speech
     : John said to Mary, "I'm dead tired and I won't go to that silly party." Voice-preserving
     narration: John was dead tired, he told Mary, and he would not go to that silly party.
SPEAK-THEN-NARRATE (improvement 2; the Actor receives the FIRST PERSON instruction, then the
     Director, temperature 0): You are the narrator of a novel. Turn the following direct speech of
      <name> (addressed to <addressees>) into third-person narration in free indirect style.
Rules: keep <name>'s own words, idioms, imagery and tone; change only what the narrative form
     requires: first-person pronouns (I, me, my, we, us, our) become <name>, he, she or they;
     second-person pronouns (you, your) become the addressee's name, he, she or they; verbs shift
     to the past tense. Do not use quotation marks or the pronouns I, me, my, we, us, our, you or
     your. Do not add commentary, interpretations or feelings that are not in the speech; at most
     add a short reporting phrase such as "he told Mary".
```

```
Example. Direct speech of John to Mary: I’m dead tired and I won’t go to that silly party.
     Narration: John was dead tired, he told Mary, and he would not go to that silly party.
Speech: <the Actor’s turn>
Return only JSON: {"content": "..."}
```